Qwen

Engineering the Uncanny: A Secure, Interoperable Framework for Quantifying and Generating Horror Art Styles

The development of a unified, secure routing framework for horror-themed AI image generation represents a significant step towards standardizing and professionalizing the creation of generative art. The core objective is to move beyond isolated, intuitive prompt engineering by establishing a systematic, data-driven methodology that unifies quantitative style metrics, measured color palettes, and decodable text-based descriptors into a single, interoperable system . This approach aims to enable consistent, cross-platform generation, analysis, and sharing of a broad taxonomy of horror artstyles, far exceeding initial examples like TerraScape and Mushroom-Drips . The framework is designed to encapsulate all relevant artifacts—including seeds, stylemaps, hex codes, and art-strings—in a secure, decodable format using Lua, thereby facilitating a robust ecosystem for creators, researchers, and developers . This report provides a comprehensive deep-dive analysis into the proposed framework, detailing its scientific methodology, architectural blueprint, implementation challenges, and strategic implications for the future of AI-assisted artistic creation.

A Scientific Methodology for Style Engineering and Validation

The foundational principle underpinning the proposed framework is the transformation of style creation from an intuitive art form into a rigorous scientific discipline. This requires a shift away from purely textual descriptions and toward a repeatable experimental loop grounded in empirical data derived from real-world visual references . The methodology is not a one-time setup but a continuous cycle of define, generate, analyze, refine, and log, ensuring that every stylistic element is anchored in both human perceptual judgment and machine-executable parameters . The first critical action is to ground each artstyle's definition in tangible visual examples rather than abstract assumptions . For instance, instead of a theoretical description of "cosmic decay," the process begins by curating a dataset of images that exemplify this aesthetic. This curated

collection becomes the primary source material for all subsequent quantification steps. To operationalize subjective qualities like "waxiness" or "void proximity," a labeling pass is conducted where annotators use simple sliders to assign numerical values based on the reference images . These labels effectively translate human intuition into approximate quantitative targets, such as a "WaxMelt Presence" score or a "VoidProximity Score," which become the basis for later metric definitions . This approach aligns with principles from psychophysics, which seek to quantify mental perceptions into measurable physical responses, providing a scientific bridge between the artist's intent and the model's execution [51] [85].

Once a style is defined by its reference set, the next stage is to derive its quantitative anchors: style metrics and hex palettes . Style metrics represent the intensity of specific stylistic elements. While established acronyms like WMP (Wax-Melt Presence), NSG (Neon Spore Glow), and TDI (Texture Density Index) may serve as starting points, their validity must be empirically tested against the curated reference images . The key is to validate whether a high WMP value consistently produces visible dripping textures in the generated output. If not, the metric itself may need refinement, or its mapping to model controls must be adjusted. Over time, this validation loop allows for the elimination of redundant or ineffective metrics, leading to a more precise and powerful set of tools for controlling the generation process . Similarly, color consistency is addressed by treating hex palettes as measured quantities, not creative guesses . The recommended procedure involves sampling colors directly from the reference images, focusing on key regions like the sky, shadows, light sources, and subjects. These sampled colors are then clustered into functional roles—such as background base, secondary wash, highlight, and accent—and representative hex values are calculated for each cluster to form a "canonical" palette for the style . This creates a standardized color anchor that can be reliably referenced and applied. The entire process is iterative; candidate hexes are tested in controlled grids where only the palette changes, allowing for the identification of combinations that most reliably evoke the target mood . This results in a documented set of primary and secondary hex values for each style, forming the basis of a broader horror-only palette library [64] [111].

With quantitative anchors in place, the focus shifts to building and refining the "style line"—a short, minimal string of tokens that drives the visual output . Research in prompt engineering emphasizes the importance of conciseness to avoid "style soup," a phenomenon where too many conflicting modifiers cause the model to produce incoherent or diluted results . A well-crafted style line might contain roughly five to seven key phrases, such as "hand-painted dark cartoon, wax-drip textures, neon mold glow, flooded urban street, eerie silence" for the Mushroom-Drips style . To refine these lines, controlled experiments are essential. This involves generating multiple images for a fixed

subject and quality tags, varying only the style line. By systematically rating the outputs for "on-style" versus "off-style" coherence, weak or ineffective tokens can be identified and dropped, while those that clearly drive texture and mood can be promoted and strengthened . This iterative testing process builds a per-style style line that becomes the engine for driving the underlying metrics and producing consistent results . The table below outlines a conceptual structure for this experimental validation process.

Experiment ID	Base Subject	Quality Tags	Style Line	Metric Profile (WMP, NSG, etc.)	Hex Palette	Generated Seed(s)	On-Style Rating (1-10)	Notes / Observations
HRS-001	Terrifying landscape	masterpiece, highly detailed	cosmic void, crumbling planet, jaundiced moon	WMP: 0.1, NSG: 0.2, TDI: 0.4	[Hex1, Hex2, ...]	10543, 88201	8	Good void presence, but lacks expected texture density.
HRS-002	Terrifying landscape	masterpiece, highly detailed	cosmic void, crumbling planet, jaundiced moon, cracked earth texture	WMP: 0.1, NSG: 0.2, TDI: 0.8	[Hex1, Hex2, ...]	10543, 88201	9	Significantly improved texture detail matches reference.
HRS-003	Terrifying landscape	masterpiece, highly detailed	analog TV static, VHS distortion, low resolution	WMP: 0.0, NSG: 0.0, TDI: 0.6	[Analog_Palette_Hexes]	10543, 88201	7	Strong analog feel, but the static feels noisy rather than atmospheric.

The final and most critical component of the scientific loop is closing it with interrogation. After generating a batch of images intended to represent a specific style and metric profile, a sample of the outputs is fed back into an interrogator model, such as CLIP-interrogator . This model generates textual descriptions of the image, providing an objective measure of how closely the output aligns with the intended concept . For a TerraScape batch, the ideal interrogator output would include words like "crumbling planet," "void," "tree," and "jaundiced moon." If the output is instead described as "melting faces" and "neon," it indicates a failure of the current style line or metric-to-token mapping, and adjustments are needed . This feedback mechanism also applies directly to color. By extracting the dominant palette from a generated image and comparing it to the target hex set, one can measure "palette drift"—the tendency of a style to deviate from its intended color scheme over time or across different seeds [70] [71] . This provides a concrete signal to drive iteration, allowing for the adjustment of either the descriptive terms or the underlying hex code assignments to correct the drift . This continuous cycle of

experimentation and measurement ensures that style definitions remain firmly grounded in real visual examples and the actual behavior of the diffusion model, rather than abstract assumptions [84].

Architectural Blueprint for a Unified Routing Framework

The ultimate ambition of this research is to construct a unified routing framework that transcends individual styles and platforms, enabling any system to safely gather, view, or reuse generated assets . The proposed architecture is built upon a central concept: the use of Lua as a lightweight, embeddable scripting language to encapsulate and serialize all style-related artifacts into a single, decodable object . This choice is strategic, offering a balance of performance, flexibility, and ease of integration that is well-suited for a system requiring both interoperability and extensibility. The core idea is to create a standardized Lua table that serves as a self-contained specification for a given piece of AI-generated horror art or a specific artstyle. This table would contain keys for all necessary components, including the numeric metric profiles, the canonical hex palettes, the refined style line tokens, and any other relevant metadata such as the seed or generation parameters . By adopting this structured approach, the framework establishes a universal data format. Any platform or tool capable of parsing this Lua object can programmatically understand the complete recipe for a style, thus enabling consistent generation and sharing across disparate systems .

The modularity of the Lua table structure is a key advantage for long-term maintenance and scalability. As new methods for style definition emerge or as the needs of the community evolve, new keys can be added to the existing table without disrupting legacy systems. For example, if a new metric called "Atmospheric Fog Density" (AFD) is developed, it can simply be added as a new key-value pair to the style's Lua object: `style_spec["AFD"] = 0.65`. Systems that are unaware of the AFD metric can continue to function correctly, ignoring the unknown key, while newer systems can leverage it for enhanced control. This contrasts sharply with rigid, flat-file formats that require complex versioning schemes to manage changes. Furthermore, embedding all related data for a single artifact within one file prevents the common problem of data fragmentation, where seeds, prompts, and palettes are stored in separate files that can become misaligned over time . This centralized data container ensures integrity and simplifies the management of complex generation histories. The design is inspired by

modular systems seen in other domains, such as human creative workflows, which allow complex concepts to carry over across different modalities and applications [2] .

A critical aspect of this architectural blueprint is the explicit mapping of style metrics to prompt levers, which is the mechanism through which the abstract numbers in the Lua table are translated into visual features . For diffusion models that support weighted terms in their prompting syntax, this mapping is direct and powerful . A high "VoidProximity Score" (e.g., 0.9) from the Lua object could be translated into a strong positive weight for the term "endless dark void" (e.g., "(endless dark void:1.3)") and a corresponding negative weight for contradictory concepts like "crowded city" . Similarly, a high "NeonSpores Glow" value would drive strong positive weights for terms like "neon glow, fluorescent green, phosphorescent highlights" and negative weights for "natural daylight, pastel tones" . This translation logic forms the heart of the routing framework. It acts as a translator, converting the structured, quantitative data from the Lua object into the free-form, qualitative language that the diffusion model understands. The effectiveness of this mapping is not static; it must be discovered through systematic A/B testing where specific metric-driven tokens are incrementally increased or decreased while all other variables are held constant . Over multiple rounds of testing, a robust mapping is established, allowing the framework to reliably translate a user's selection of a style from the database into a well-formed prompt that generates the desired result .

The routing framework's ability to handle both positive and negative prompts is crucial for achieving stylistic purity and preventing unwanted artifacts. Negative prompts are particularly effective for enforcing stylistic boundaries. For example, a high "SpecterObfuscation Metric" might necessitate explicitly weighting anatomical clarity negatively to ensure monsters remain ambiguous silhouettes . The framework's architecture must therefore support the generation of sophisticated negative prompts based on the metric profile. For instance, a style characterized by high "Urban Decay" might trigger negative weights for "pristine surfaces, clean architecture, sunny weather." This dual-control system (positive and negative weights) allows for fine-grained sculpting of the output space. The process of discovering these mappings is part of the ongoing experiment-and-measure loop. If a particular style frequently suffers from unrealistic lighting despite having the correct texture metrics, the team might discover that adding a negative prompt tag like "volumetric lighting, studio lights" improves consistency. This discovery is then codified into the routing logic, further refining the framework's intelligence. The entire process underscores the necessity of model-specific experimentation, as the optimal weights and tokens for one model (e.g., Stable Diffusion v1.5) may not be the same for another (e.g., SDXL 1.0) [67] [101]. Therefore, the framework must be adaptable, with its internal mapping tables potentially needing to be re-tuned as models evolve.

Implementing Cross-Platform Consistency and Hex Palettes

Achieving consistent style application across diverse AI image generation platforms is a central challenge for the proposed framework. Different platforms offer varying degrees of control over the generation process, making a one-size-fits-all approach untenable. The strategy hinges on a dual-method approach for handling hex palettes and style descriptors, accommodating both direct color control where available and text-based approximations backed by reference images for more limited interfaces . This ensures that the core visual identity of a horror style, defined by its quantitative hex palette, can be preserved regardless of the target platform's capabilities . Some modern generation tools, such as Amazon's Titan Image Generator G1 models, provide explicit API endpoints that allow users to supply a list of hex color codes alongside a text prompt to guide the image generation process [39]. When interfacing with such a platform, the routing framework can directly extract the canonical hex palette from the Lua object and populate the appropriate API call, ensuring the highest possible fidelity to the intended color scheme. This direct method is the gold standard for color control and should be prioritized whenever the target platform supports it.

However, a significant number of platforms do not offer direct hex input. For these environments, the framework must rely on a combination of descriptive text and contextual guidance . The process involves translating the canonical hex codes into evocative, human-readable color names that capture their perceived quality. For example, a dark, murky blue-green hex code might be labeled "deep sewer green" or "foggy river black," while a toxic-looking yellow-green could be "toxic neon green" [64]. These descriptive tags are then woven into the style line, providing the model with semantic cues about the desired color palette. To further anchor these text-based descriptions, the framework leverages the power of reference images. During generation, a small thumbnail of the reference image that originally yielded the canonical palette is passed to the model along with the prompt . This multi-modal input provides the model with a visual anchor, helping it interpret the descriptive text in the context of the correct hues. The effectiveness of this method must be rigorously tested. Controlled experiments should compare outputs generated with and without the reference image, and with different sets of descriptive tags, to determine which combinations most reliably reproduce the target palette on a given platform . This empirical validation is key to building a robust library of both primary hex codes and their corresponding descriptive equivalents.

The process of deriving and validating these hex palettes is a meticulous one. It begins with sampling colors from a carefully curated set of reference artworks that exemplify the target horror style . Advanced techniques like color quantization using K-Means Clustering can be employed via libraries such as OpenCV to algorithmically reduce the color palette of an image and identify its dominant hues [31] [71]. Once a set of candidate colors is extracted, they are clustered into distinct roles: a base/background color, secondary washes, highlights, and accents . Representative hex values are then calculated for each cluster, forming the "canonical" palette for the style. This palette is not a suggestion but a quantitative target. The framework's logging and analysis loop is instrumental here. By periodically re-sampling the hex palettes from successfully generated images, the system can track for "palette drift"—where the model's output gradually shifts away from the original target colors [70]. If a particular style consistently generates images that are too saturated or have an incorrect hue balance, the canonical hex set can be updated, and the descriptive tags can be refined to better guide the model . Over time, this process yields a comprehensive horror palette atlas, where each style has a validated set of tested hex clusters [64]. This atlas can be used to create crossovers or new variants while maintaining tight color discipline, ensuring that even novel combinations adhere to a coherent visual language .

To manage this complexity, a structured approach is necessary. The following table illustrates how a single horror style, "Wax-Melt Urban Horror," could be defined with its corresponding hex palette and descriptive tags for use in the framework.

Palette Role	Canonical Hex Code	Description for UI/ API	Descriptive Tag for Prompt	Example Use in Prompt
Primary Base	`#1A1E26`	Deep Asphalt Blue	"dark asphalt streets, foggy pavement"	"flooded urban street at night, dark asphalt streets..."
Secondary Wash	`#2C313C`	Shadowed Concrete Grey	"shadowed concrete walls, rainy reflections"	"...shadowed concrete walls covered in..."
Highlight	`#FF3E00`	Toxic Neon Orange	"sodium-vapor orange light, glowing puddles"	"pools of toxic neon orange liquid reflecting..."
Accent	`#00FF88`	Phosphorescent Green	"neon mold glow, fluorescent green spores"	"neon mold glow emanating from the..."
Skin/Wax Tone	`#F4E3B9`	Jaundiced Moon Yellow	"jaundiced moon-yellow wax, melting skin"	"figures with jaundiced moon-yellow wax..."

This structured mapping ensures that whether a platform accepts a hex code directly or relies on text, the core color identity of the style is maintained. The framework's adaptability is its greatest strength, allowing it to function effectively across a fragmented ecosystem of AI art generators. The continuous validation loop, combining direct hex

input where possible and rigorous testing of text+image approximations elsewhere, is what ultimately guarantees cross-platform consistency and fidelity [39].

The Role of Decodable Artifacts and Data Logging

A cornerstone of the proposed framework is the comprehensive management and logging of all artifacts generated during the style creation process . This includes not only the final images but also the seeds, stylemaps, hex codes, and art-strings that produced them. By encoding these artifacts into a decodable format, primarily using Lua, the system transforms a series of isolated creations into a structured, queryable dataset . This approach moves beyond simple asset storage and positions the framework as a dynamic research tool for continuous improvement. Every successful or unsuccessful generation run is meticulously recorded, creating a rich historical dataset that captures the relationship between input parameters and visual outcomes . This log becomes the foundation for a formal, data-driven refinement process, allowing developers to analyze patterns, diagnose failures, and enhance the precision of their style definitions over time .

The act of encoding all artifacts into a single Lua object is central to this strategy. This decodable value serves as a complete snapshot of a generation event . It contains the full specification of the style being tested, including the numeric values for all relevant metrics (e.g., WMP, NSG, TDI), the exact hex codes for the palette, the final, optimized style line, and the specific seed or range of seeds used. This level of detail is invaluable. If a particular seed produces a stunningly good or catastrophically bad result, the log allows for a full forensic reconstruction of the conditions that led to that outcome. This is far more informative than simply saving an image file. It enables a deeper understanding of model behavior, revealing which seeds are consistently "good" for a given style and which are "wildcard" seeds that produce unpredictable results. The ability to regenerate an image exactly as it was produced is a powerful feature for both quality assurance and collaborative work, ensuring that a discovered "perfect" seed can be reliably reused [53] [80].

The logging system must be designed to capture a wide range of metadata for each experiment. Beyond the parameters explicitly defined in the Lua object, the log should record the timestamp, the user or agent responsible, the specific version of the diffusion model used, and the hardware environment. Crucially, it must also store the output of the interrogator model that runs after generation . This interrogator summary provides an objective textual description of the generated image, which can be compared against the intended description to calculate a "style alignment score" [42]. This automated

evaluation metric is a key driver for the refinement loop. If a large number of logs show that a certain style consistently receives interrogator tags like "realistic" when the goal was "stylized cartoon," it signals a clear direction for adjustment, perhaps by modifying negative prompt weights or refining the style line itself . This turns the manual review process into a scalable, data-centric workflow.

Over time, the accumulated log data becomes a powerful analytical resource. It allows for statistical analysis of the correlation between different metrics and visual features. For example, by querying the database, one might discover that a high "Texture Density Index" (TDI) almost always corresponds to more visible dripping textures, confirming the metric's validity. Conversely, another metric might show no significant correlation with any observable feature, indicating it is redundant and should be deprecated . This data-driven approach to metric validation is a significant departure from relying on anecdotal evidence. It also enables the discovery of emergent properties and failure modes that were not initially anticipated. For instance, the logs might reveal that pushing the "VoidProximity" metric too high on a particular model causes it to generate nonsensical floating objects, a failure mode that can then be specifically targeted with new negative prompts. This continuous analysis loop aligns with methodologies used in advanced AI safety research, such as automatic red-teaming frameworks (ART), which use systematic evaluation and logging to uncover vulnerabilities in text-to-image models [3] . By applying similar principles to the domain of artistic style, the framework ensures that its definitions are not static but are constantly evolving and improving based on empirical evidence from the generation process itself [3] [83] .

Security Architecture for a Safe and Trustworthy System

The explicit requirement for a secure framework to share sensitive data like generation seeds introduces a critical layer of complexity that must be addressed at the architectural level . While the ability to share seeds, stylemaps, and other artifacts is a core tenet of the framework's utility, it also presents a potential attack vector. Malicious actors could craft seemingly benign Lua objects containing harmful code that, upon execution, could compromise a user's system or data. Therefore, a robust security architecture is not an optional add-on but a fundamental requirement for the framework's safe operation. This architecture must be designed to prevent the execution of arbitrary code while still allowing for the interpretation of the Lua object as a data structure. The security model should be predicated on the principle of least privilege, ensuring that the process of

decoding and using the Lua object has access only to the absolute minimum resources required for image generation.

One of the primary security risks associated with programmable languages like Lua is code injection [92]. A malicious actor could attempt to embed executable Lua functions within the data fields of a style specification, hoping to trick a vulnerable parser into running them. Historical vulnerabilities in software have shown that local attackers can bypass security restrictions by placing specially crafted Lua bytecode files in configuration directories [92]. To counter this, the framework must implement a "safe" or "sandboxed" Lua environment. Creating a secure Lua sandbox involves restricting the interpreter's access to dangerous functions and APIs [90]. In Lua 5.2 and later, functions like `load` and `loadfile` can be encapsulated to limit their capabilities, for example, by forcing a specific mode parameter that prevents execution [90]. More broadly, the sandbox should disable or override standard library functions that provide access to the filesystem, network, operating system commands, or other external processes. This ensures that even if a malicious payload is embedded within the Lua data, it cannot break out of the controlled environment to cause harm. The security of microelectronics and software components is a growing concern, and such strict scrutiny is necessary to mitigate emerging cyber threats [8].

In addition to sandboxing the interpreter, the framework must employ strict data validation and schema enforcement before any data is processed. All incoming Lua objects should be validated against a predefined, whitelisted schema that specifies the acceptable keys and data types for each field. For example, the schema would declare that `"WMP"` must be a floating-point number between 0 and 1, `"hex_palette"` must be a table of valid hex string representations, and the `"style_line"` must be a string. Any key or value that does not conform to this schema should be rejected or ignored, preventing unexpected or malicious data from influencing the generation process. This practice is analogous to securing AI agents against prompt-injection attacks, where malformed inputs are used to subvert the agent's intended behavior [5] [7]. The use of secure coding practices and regular security audits of the underlying libraries, such as Python packages, is also essential to patch known vulnerabilities [12]. Apache Airflow, for instance, disabled deserialization of arbitrary XCom values in its API endpoints for security reasons, a practice that underscores the importance of not blindly trusting serialized data [9].

Finally, the framework should enforce immutability for the decoded artifacts. Once a Lua object is parsed and validated, the resulting configuration data should be treated as read-only for the duration of the generation task. This prevents any possibility of the

generation process or the model itself altering the original style specification. This principle of immutability is a key tenet of secure systems design, as it ensures that the input remains consistent and untampered with throughout the execution lifecycle. The entire security architecture—from sandboxing and data validation to immutability—must be designed to protect the user's system while still enabling the core functionality of securely sharing and reusing the rich, decodable style specifications. Without this robust security posture, the promise of a collaborative and interoperable horror art ecosystem would be undermined by significant risk.

Strategic Synthesis: From Prompt Engineering to a Taxonomy of Horror

The development of this unified routing framework marks a strategic evolution from ad-hoc prompt engineering to the establishment of a formal, extensible taxonomy of horror artstyles. While early AI art creation often relied on accumulating lists of "art styles" and "prompt modifiers" [34], this project seeks to build a more profound and durable system. The framework's true innovation lies in its emphasis on the relationships between styles, not just their individual definitions. By creating a "horror style matrix" that compares different archetypes—such as TerraScape, Mushroom-Drips, analog ghoul silhouettes, and others—it becomes possible to map the boundaries and overlaps between them . This comparative analysis is vital for preventing "style drift," a common problem where pushing certain parameters too far inadvertently pushes a generation into the territory of a different, unintended style . For example, excessive use of neon effects in a cosmic decay scene might unintentionally make it look like a Mushroom-Drips urban horror piece unless urban artifacts are specifically constrained .

This comparative approach leads to the development of a hierarchical vocabulary for horror aesthetics. At the top level, the framework identifies "horror-global" tokens that establish a baseline mood applicable to the genre as a whole. These might include words like "eerie," "unsettling," "foreboding," and "cinematic horror lighting" . These global tokens provide a consistent atmospheric foundation. Below this, the framework defines style-specific tokens that give each artstyle its unique identity. For "cosmic void," this might be "crumbling planet," "jaundiced moon," and "infinite darkness," while for "analog ghoul silhouettes," it could be "VHS static," "low-resolution distortion," and "CRT screen glow" [116]. The routing logic can then be designed to stack these layers, combining the global horror mood with the specific identity of a chosen artstyle. This modular approach allows for both consistency and variety, enabling the creation of novel combinations

while maintaining the integrity of the core styles. The research involved in this taxonomy development focuses on determining which global tokens are truly safe to reuse across styles and which ones tend to pull disparate styles too close together, erasing their unique characteristics .

Ultimately, the framework is a system designed for continuous learning and adaptation. Its reliance on a perpetual experiment-and-measure loop ensures that style definitions do not become static dogma but living documents that evolve with the technology . The rapid pace of development in diffusion models means that a perfectly tuned style definition today may degrade in quality tomorrow [67] . The framework's logging and analysis capabilities are the antidote to this problem. By continuously monitoring the performance of each style, the system can flag degradation and trigger a re-validation process. This proactive maintenance is crucial for the long-term viability of the taxonomy. The uncanny valley, a psychological concept describing the unsettling feeling caused by near-human replicas, is a fitting metaphor for the challenges of horror aesthetics [20] [24] . The framework's granular control over texture, lighting, and color is precisely the kind of tool needed to navigate this space, deliberately crafting visuals that are stylized enough to be compelling but realistic enough to be disturbing. The success of the framework will depend not only on its technical implementation but also on its ability to foster a community of practice around this new, data-driven methodology for creating AI-assisted art.

Reference

1. Civiverse: A Dataset for Analyzing User Engagement with Open ... https://arxiv.org/html/2408.15261v1
2. Ontology-Driven Multi-Agent System for Cross-Domain Art Translation https://www.mdpi.com/1999-5903/17/11/517
3. (PDF) ART: Automatic Red-teaming for Text-to-Image Models to ... https://www.researchgate.net/publication/381006713_ART_Automatic_Red-teaming_for_Text-to-Image_Models_to_Protect_Benign_Users
4. In search of indicators of culture and development: progress and ... https://unesdoc.unesco.org/ark:/48223/pf0000121133
5. [PDF] AI Security in the Foundation Model Era: A Comprehensive Survey ... https://openreview.net/pdf?id=1g7pKgClZs

6. Arxiv今日论文| 2026-03-17 - 闲记算法 http://lonepatient.top/2026/03/17/arxiv_papers_2026-03-17.html
7. Building AI Coding Agents for the Terminal: Scaffolding, Harness ... https://arxiv.org/html/2603.05344v1
8. Secure and Trustworthy Microelectronics: Vulnerabilities, Solutions ... https://link.springer.com/article/10.1007/s41635-025-00165-x
9. Release Notes - Apache Airflow https://airflow.apache.org/docs/apache-airflow/2.8.4/release_notes.html
10. Machine Learning in Python: Main Developments and Technology ... https://www.mdpi.com/2078-2489/11/4/193
11. Exploitable C# Functions - security - Stack Overflow https://stackoverflow.com/questions/3940576/exploitable-c-sharp-functions
12. Security update for python-arcomplete, python-Fabric,... - SUSE https://www.suse.com/zh-cn/support/update/announcement/2024/suse-su-20241639-1/
13. TARmageddon: A Critical Rust Vulnerability for Kubernetes and AI https://www.linkedin.com/posts/deepthi-talasila_how-tarmageddon-compromises-rust-security-activity-7392349068162842624-6ybl
14. Do You See Me : A Multidimensional Benchmark for Evaluating ... https://arxiv.org/html/2506.02022v1
15. [PDF] benchmarking large vision-language models - OpenReview https://openreview.net/pdf/da29d197514ee7c04c2724a3ae8cd668ed02f510.pdf
16. [PDF] Large Vision-Language Model Alignment and Misalignment https://aclanthology.org/anthology-files/anthology-files/pdf/findings/2025.findings-emnlp.90.pdf
17. STATUS Bench: A Rigorous Benchmark for Evaluating Object State ... https://dl.acm.org/doi/10.1145/3746027.3755565
18. A Survey of State of the Art Large Vision Language Models https://arxiv.org/html/2501.02189v4
19. A Survey of State of the Art Large Vision Language Models https://www.researchgate.net/publication/390334774_A_Survey_of_State_of_the_Art_Large_Vision_Language_Models_Alignment_Benchmark_Evaluations_and_Challenges
20. Uncanny as Antisublime | The Journal of Aesthetics and Art Criticism https://academic.oup.com/jaac/article/83/4/330/8257125
21. The Uncanny City and the Immanence of Crisis: An Arts-Informed ... https://link.springer.com/article/10.1007/s42087-025-00483-0

22. Iwaa and Sophia robot versus a real human being - ScienceDirect https://www.sciencedirect.com/science/article/pii/S1071581923001519
23. audiovisual and psychological dimensions of creepypasta https://www.researchgate.net/publication/399325920_AUDIOVISUAL_AND_PSYCHOLOGICAL_DIMENSIONS_OF_CREEPYPASTA_A_COMPARATIVE_MEDIA_ANALYSIS
24. (PDF) Perception of psychopathy and the Uncanny Valley in virtual ... https://www.academia.edu/31958556/Perception_of_psychopathy_and_the_Uncanny_Valley_in_virtual_characters
25. Fear in Music – Expression, Perception, and How It Actually Feels https://journals.sagepub.com/doi/10.1177/20592043261432393
26. 实用英汉计算机词汇大全 | PDF - Scribd https://www.scribd.com/document/896223284/%E5%AE%9E%E7%94%A8%E8%8B%B1%E6%B1%89%E8%AE%A1%E7%AE%97%E6%9C%BA%E8%AF%8D%E6%B1%87%E5%A4%A7%E5%85%A8
27. 333333 23135851162 the 13151942776 of 12997637966 ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt
28. Google 1-gram - CMU School of Computer Science https://www.cs.cmu.edu/~roni/11761/2017_fall_assignments/hw3_stats_google_1gram.txt
29. Spam-Ham NLP Basic Project - Kaggle https://www.kaggle.com/code/syamalakumar/spam-ham-nlp-basic-project
30. Multifaceted Gold–Palladium Bimetallic Nanorods and Their ... https://pubs.acs.org/doi/10.1021/acsnano.7b00264
31. How to reduce the number of colors in an image with OpenCV? https://stackoverflow.com/questions/5906693/how-to-reduce-the-number-of-colors-in-an-image-with-opencv
32. Hex-Mesh Generation and Processing: A Survey - ACM Digital Library https://dl.acm.org/doi/full/10.1145/3554920
33. Ultimate guide to PyTorch library in Python - Deepnote https://deepnote.com/blog/ultimate-guide-to-pytorch-library-in-python
34. [PDF] Using Multimodal Interaction for Design Exploration with Generative AI https://hal.science/hal-04652020v1/file/DIS_24_DesignPrompt_Author_Version.pdf
35. Using Multimodal Interaction for Design Exploration with Generative AI https://www.researchgate.net/publication/381743475_DesignPrompt_Using_Multimodal_Interaction_for_Design_Exploration_with_Generative_AI
36. Anime Generation through Diffusion and Language Models https://www.sciencedirect.com/org/science/article/pii/S1526149225002991

37. [PDF] Protocol-level Attacks and Defenses to Advance IoT Security https://theses.hal.science/tel-05019219v1/file/145708_CASAGRANDE_2024_archivage.pdf
38. Findings of the Association for Computational Linguistics (2024) https://aclanthology.org/events/findings-2024/
39. Amazon Titan Image Generator G1 models - Amazon Bedrock https://docs.aws.amazon.com/bedrock/latest/userguide/model-parameters-titan-image.html
40. [PDF] arXiv:2503.23907v1 [cs.CV] 31 Mar 2025 https://arxiv.org/pdf/2503.23907
41. Perceptual Visual Quality Assessment: Principles, Methods, and ... https://www.researchgate.net/publication/389548640_Perceptual_Visual_Quality_Assessment_Principles_Methods_and_Future_Directions
42. [PDF] EMONET-FACE: An Expert-Annotated Benchmark for ... - OpenReview https://openreview.net/pdf/ac4fb611377f419b4f3a00082d745d7d353b230e.pdf
43. "Special Relativity" of Image Aesthetics Assessment: a Preliminary ... https://dl.acm.org/doi/10.1145/3664647.3681172
44. AI-Powered Procedural Haptics for Narrative VR - MDPI https://www.mdpi.com/2414-4088/10/1/9
45. Text − to − Haptics: Enhancing Multisensory Storytelling through ... https://advanced.onlinelibrary.wiley.com/doi/10.1002/aisy.202400758
46. The Uncanny Child in Transnational Cinema - jstor https://www.jstor.org/content/pdf/oa_book_monograph/10.2307/j.ctv80cc7v.pdf
47. A study of the evaluation metrics for generative images ... https://www.researchgate.net/publication/369493032_A_study_of_the_evaluation_metrics_for_generative_images_containing_combinational_creativity
48. Best Free AI Tools for Image Generation & Editing in 2026 - Toolify https://www.toolify.ai/free-ai-tools/image-generation-editing
49. SynthID-Image: Image watermarking at internet scale - arXiv https://arxiv.org/html/2510.09263v1
50. [XML] https://stackoverflow.com/sitemap-questions-271.xml https://stackoverflow.com/sitemap-questions-271.xml
51. Contrast Sensitivity in Multimodal Large Language Models https://www.sciencedirect.com/science/article/pii/S0893608026003643
52. bing.txt - FTP Directory Listing ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/bing.txt
53. SeedEdit 3.0: Fast and High-Quality Generative Image Editing - arXiv https://arxiv.org/html/2506.05083v2

54. Benchmarking autoregressive conditional diffusion models for ... https://www.sciencedirect.com/science/article/pii/S0893608026001036
55. Parameter Efficient Tuning of Text-to-Image Diffusion Models https://dl.acm.org/doi/full/10.1145/3730403
56. Sample-Adaptive Reconstruction Via Latent Diffusion Models - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC11421836/
57. Diffusion-Based Anonymization and Foundation Model-Powered ... https://www.mdpi.com/2032-6653/17/2/70
58. Training-Free Distribution Adaptation for Diffusion Models via ... https://arxiv.org/html/2601.08379v1
59. Dynamic community detection using class preserving time series ... https://www.nature.com/articles/s41598-026-37699-1
60. 333333 23135851162 the 13151942776 of 12997637966 https://www.cs.princeton.edu/courses/archive/spring18/cos226/assignments/autocomplete/testing/words-333333.txt
61. cmnt_vocab.txt - CMU School of Computer Science https://www.cs.cmu.edu/~ark/blog-data/data/blog_data_v1_0/dk/hbc_data/data/cmnt_vocab.txt
62. (PDF) Processing a handbook - Academia.edu https://www.academia.edu/37221147/Processing_a_handbook
63. Evolution of American Poetry | PDF - Scribd https://www.scribd.com/document/837967576/The-Cambridge-History-of-American-Poetry
64. English.txt - APRIL robotics lab https://april.eecs.umich.edu/courses/eecs492_w11/wiki/images/8/82/English.txt
65. [PDF] ARCS Workshops - Braunschweig, Germany - ResearchGate https://www.researchgate.net/profile/Onur-Kilincceker-2/publication/329390717_Applying_the_Ideal_Testing_Framework_to_HDL_Programs/links/61caeda9e669ee0f5c6c00d0/Applying-the-Ideal-Testing-Framework-to-HDL-Programs.pdf
66. Collaborate Computing: Networking, Applications and Worksharing https://link.springer.com/content/pdf/10.1007/978-3-319-59288-6.pdf
67. Backdooring Text-to-Image Models by Masquerading as Benign ... https://arxiv.org/html/2602.21977v3
68. Personalized smart immersive XR environments: a systematic ... https://link.springer.com/article/10.1007/s00371-025-03887-9
69. Funny Addons collection (weird alert) - Steam Community https://steamcommunity.com/sharedfiles/filedetails/?id=2848029040

70. Real-time Image Color Palette Extractor — A Deep Dive into K ... https://dev.to/ertugrulmutlu/real-time-image-color-palette-extractor-a-deep-dive-into-k-means-lab-and-de2000-4eoi
71. Color Quantization with OpenCV using K-Means Clustering https://pyimagesearch.com/2014/07/07/color-quantization-opencv-using-k-means-clustering/
72. Exclusive Magazine - DVD Reviews - Anne Carlini https://annecarlini.com/ex_dvd.php
73. Why Quasar PWA + SSR server-rendered source code not work for ... https://stackoverflow.com/questions/78512608/why-quasar-pwa-ssr-server-rendered-source-code-not-work-for-production-build-g
74. A top down description of S − CIELAB and CIEDE2000 - ResearchGate https://www.researchgate.net/publication/227604556_A_top_down_description_of_S-CIELAB_and_CIEDE2000
75. A comparison of the CIELab and CIEDE2000 color difference formulas https://www.researchgate.net/publication/282759670_A_comparison_of_the_CIELab_and_CIEDE2000_color_difference_formulas
76. Color Appearance and Color Difference Specification - ResearchGate https://www.researchgate.net/publication/279426807_Color_Appearance_and_Color_Difference_Specification
77. Search the Listings - The Webcomic List http://www.thewebcomiclist.com/search.php?q=best+web+
78. ClipFaceFusion multi modal diffusion for high fidelity facial ... - Nature https://www.nature.com/articles/s41598-025-31331-4
79. [PDF] Structure and Content-Guided Video Synthesis with Diffusion Models https://openaccess.thecvf.com/content/ICCV2023/papers/Esser_Structure_and_Content-Guided_Video_Synthesis_with_Diffusion_Models_ICCV_2023_paper.pdf
80. Discovering Secret Seeds in Text-to-Image Diffusion Models - arXiv https://arxiv.org/html/2405.14828v1
81. [PDF] ConsistI2V: Enhancing Visual Consistency for Image-to-Video ... https://openreview.net/pdf/e3c4edb75a0e7089a966f4481203f9b8a4bc9e32.pdf
82. DAttnVis: Attention-guided visual diagnostics for stable diffusion ... https://journals.sagepub.com/doi/10.1177/14738716261434840
83. Discovering Secret Seeds in Text-to-Image Diffusion Models https://www.researchgate.net/publication/390617446_Good_Seed_Makes_a_Good_Crop_Discovering_Secret_Seeds_in_Text-to-Image_Diffusion_Models

84. Synthesis and perceptual scaling of high-resolution naturalistic ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12696119/
85. Contrast sensitivity function in deep networks | Request PDF https://www.researchgate.net/publication/370268032_Contrast_sensitivity_function_in_deep_networks
86. [PDF] Pictorial composition: modeling, perception & creation - HAL theses https://theses.hal.science/tel-04646624v1/file/Lelievre_2022_These.pdf
87. European Conference on Visual Perception 2024 - Sage Journals https://journals.sagepub.com/pb-assets/cmscontent/PEC/ECVP%202024%20Abstracts-1742366759.pdf
88. European Conference on Visual Perception 2017 - Sage Journals https://journals.sagepub.com/pb-assets/cmscontent/PEC/ECVP2017_abstracts_1.0.pdf
89. 38 - Housewatch https://www.housewatch.com.au/searchautocomplete/suggest/result/
90. Creating a secure Lua sandbox..? - Stack Overflow https://stackoverflow.com/questions/34388285/creating-a-secure-lua-sandbox
91. [PDF] Lmod Documentation https://lmod.readthedocs.io/_/downloads/en/latest/pdf/
92. CVSS score 8-9 vulnerabilities - CVEs - page 2 - Feedly https://feedly.com/cve/severity/8-9?page=2
93. 85 Analog horror ideas to save today - Pinterest https://www.pinterest.com/coniimations/analog-horror/
94. Discover 100 Horror references and art inspiration ideas - Pinterest https://www.pinterest.com/elsbethgamez/horror-references/
95. Horror Composition - Pinterest https://www.pinterest.com/ideas/horror-composition/902269410090/
96. 280 Analog Horror ideas in 2026 - Pinterest https://www.pinterest.com/ritaleader15/analog-horror/
97. dialogue between body and sound through multimodal interaction https://www.academia.edu/116277649/DIALOGUE_BETWEEN_BODY_AND_SOUND_THROUGH_MULTIMODAL_INTERACTION_
98. Multi-scale Latent Point Consistency Models for 3D Shape Generation https://arxiv.org/html/2412.19413v1
99. [PDF] Multi-Guidance Diffusion for Controllable Zero-Shot Immersive ... https://openaccess.thecvf.com/content/ICCV2025/papers/Hu_MiDSummer_Multi-Guidance_Diffusion_for_Controllable_Zero-Shot_Immersive_Gaussian_Splatting_Scene_ICCV_2025_paper.pdf

100. [PDF] MULTI-SCALE LATENT POINT CONSISTENCY MODELS FOR 3D ... https://openreview.net/pdf?id=92FEM1voOW
101. Quantitative comparison of fine-tuning techniques for pretrained ... https://www.sciencedirect.com/science/article/pii/S0924271626000687
102. Proceedings of the 30th ACM International Conference on Multimedia https://dl.acm.org/doi/proceedings/10.1145/3503161?tocHeading=heading13
103. Publications - APRIL Lab https://april.zju.edu.cn/publications/
104. Universal Principles of Design https://people.engr.tamu.edu/pcr/courses/csce431/winter20/UniversalPrinciplesOfDesign.pdf
105. ALIGNED BENCHMARKS IN SPEECH ROLE-PLAY https://openreview.net/pdf?id=FOZ4FwC7YX
106. Visual Representations and Interpretations - Springer Nature https://link.springer.com/content/pdf/10.1007/978-1-4471-0563-3.pdf
107. Algorithms and Architectures for Parallel Processing - Springer Nature https://link.springer.com/content/pdf/10.1007/978-981-95-8411-6.pdf
108. Online Engineering | PDF | Internet Of Things - Scribd https://www.scribd.com/document/739508613/Online-Engineering
109. dictionary.txt http://www.cs.cmu.edu/~ab/15-110F10/labs/lab4/startercode/dictionary.txt
110. Human-Computer Interaction - Springer Nature https://link.springer.com/content/pdf/10.1007/978-3-031-93982-2.pdf
111. words.txt - Apache's svn http://svn.apache.org/repos/asf/juddi/tags/juddi-3.0.1/qa/juddi-xlt/config/data/default/words.txt
112. Afrikaans English | PDF - Scribd https://www.scribd.com/document/933386190/Afrikaans-English
113. Override "user agent stylesheet" on focus - css - Stack Overflow https://stackoverflow.com/questions/53213639/override-user-agent-stylesheet-on-focus
114. Coding Style Guidelines — Carbonite SDK - NVIDIA Omniverse https://docs.omniverse.nvidia.com/kit/docs/carbonite/latest/CODING.html
115. The Weird - MCG - Core Rulebook (OEF, 2023-09-21) | PDF - Scribd https://www.scribd.com/document/701095024/The-Weird-MCG-Core-Rulebook-OEF-2023-09-21
116. wikilex-20070908-zh-en.txt - CMU School of Computer Science https://www.cs.cmu.edu/afs/cs.cmu.edu/project/cmt-40/Nice/Transfer/Chinese/wikilex-20070908-zh-en.txt
117. Challenges in Text Analysis - Fully connected NN - Kaggle https://www.kaggle.com/code/farsanas/challenges-in-text-analysis-fully-connected-nn